

Azure Content Understanding

Global AI community

12 mars 2026

Présentation du service Azure Content Understanding



**global AI
community**
Virtuelle Francophone

Présentation du service Azure Content Understanding


Azure Content Understanding est un service Microsoft Foundry qui **automatise l'extraction et la structuration de contenu multimodal (documents, images, vidéos, fichiers audios) sans labélisation complexe.**

Une présentation de ce service à l'aide de démonstrations sera réalisée lors de ce webinaire.



Serge RETKOWSKY

France AI Global Black Belt
@Microsoft



Randa MEZRAG

Consultante DATA & IA
@Finelog-Biseum

LinkedIn Live

jeudi

12/03/2026

à partir de 13h

<https://www.linkedin.com/events/7428034373116928000>

Agenda

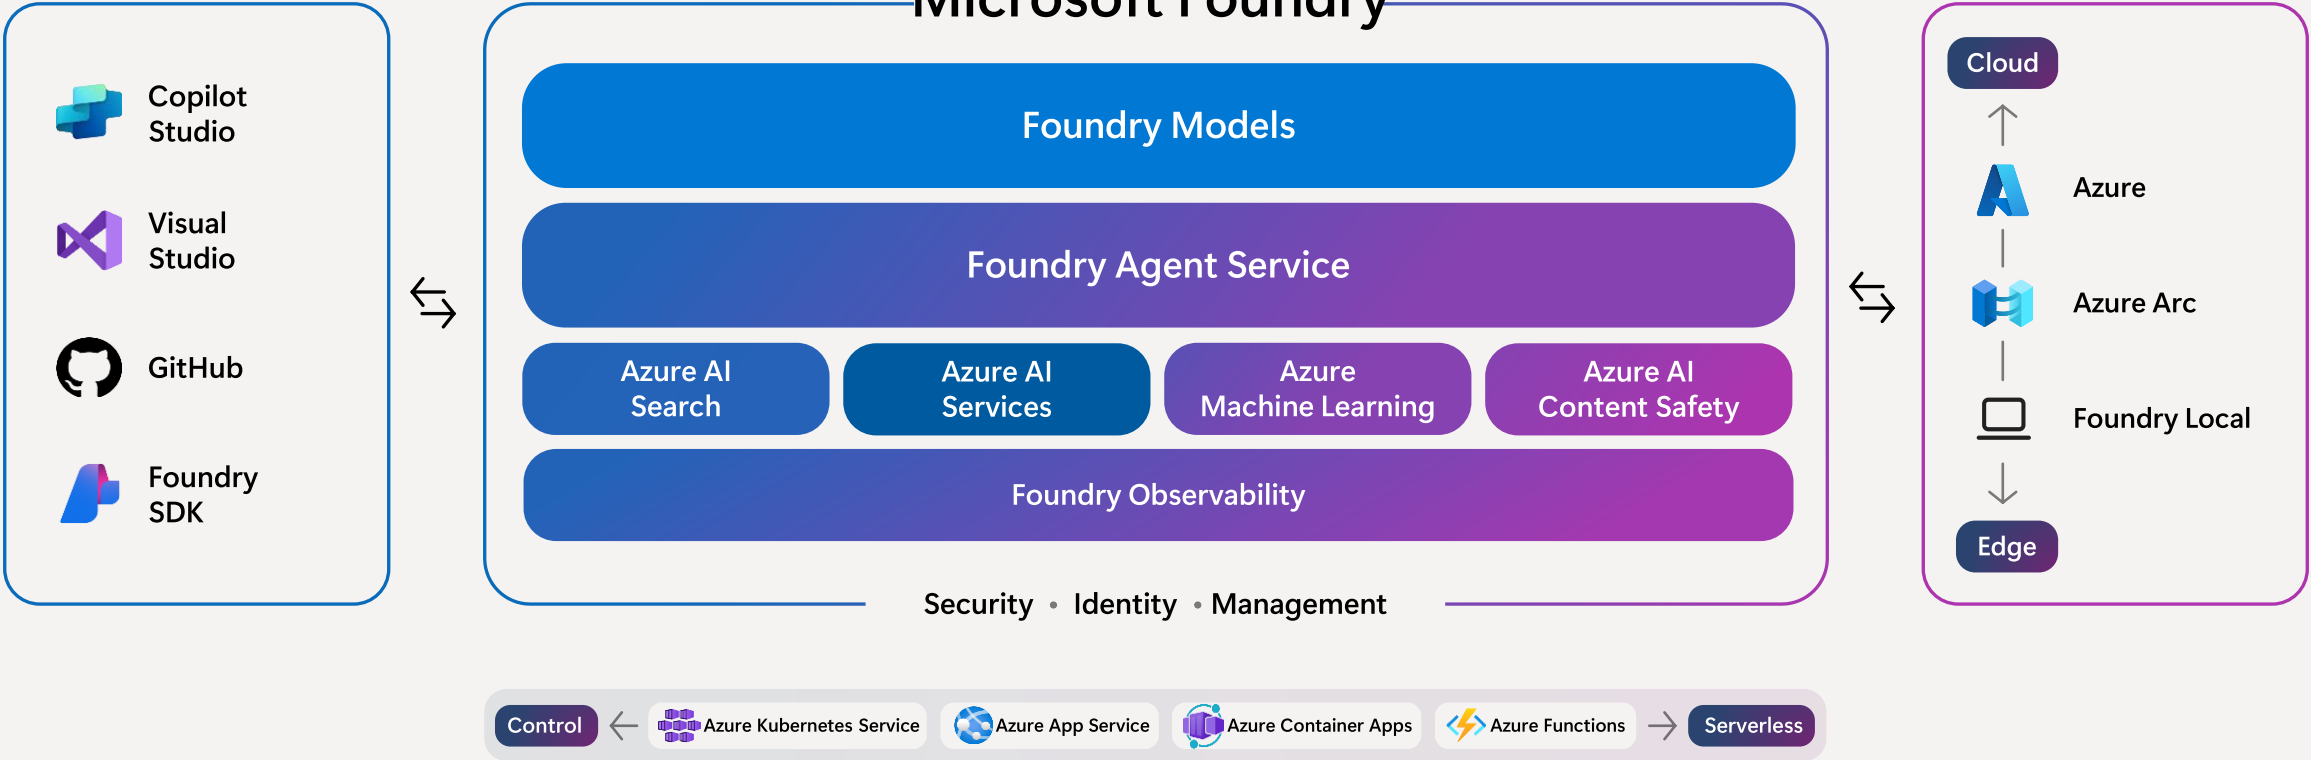
- **Présentation Azure Content Understanding**
- Démonstrations
- Azure Document Intelligence vs Azure Content Understanding
- Ressources et documentation



Microsoft Foundry



Microsoft Foundry



Foundry Tools – Azure services

Content Understanding



- Multi-file processing
- Multi-modal processing
- Prebuilt Analyzers
- Documents
- Images
- Video
- Audio

Document Intelligence



- Document Analysis
- Prebuilt Models
- Custom
- Add-on Capabilities

Language



- Entity extraction (PII/PHI)
- Summarization
- Text analytics
- Sentiment mining
- Conversational AI

Speech



- Speech to text
- Text to speech
- Custom Neural Voice
- Speech Translation
- Avatar
- Embedded speech

Translator



- Text translation
- Document translation
- Speech translation
- Translator Pro
- Embedded translation

Vision



- OCR
- Facial Recognition

Foundry Tools Availability Options



Microsoft
Public Cloud



Microsoft
Government Cloud



Connected
Containers



Disconnected
Containers


Regulatory
Compliance



HITRUST



Content Understanding framework



Content Understanding is now a Generally Available (GA) service with the release of the 2025-11-01 API version.

Content Understanding

Microsoft Foundry / demos-resource-project

Search with AI (Ctrl + K)

New Foundry

Home

Discover

Build

Operate

Docs

Agents

Workflows

Models

Fine-tune

Tools

Knowledge

Data

Evaluations

Guardrails

Models

Deployments AI Services Batch Jobs

Name	Version	Type
☐ Azure Speech - Voice Live	N/A	Speech
☐ Azure Speech - Speech to Text	N/A	Speech
☐ Azure Speech - Text to Speech	N/A	Speech
☐ Azure Speech - Text to Speech Avatar	N/A	Speech
☐ Azure Translator - Text translation	N/A	Translation
☐ Azure Translator - Document translation	N/A	Translation
☐ Azure Language - Language detection	N/A	Language
☐ Azure Language - Text PII redaction	N/A	Language
☐ Azure Content Understanding - Read	N/A	Content Understanding
☐ Azure Content Understanding - Layout	N/A	Content Understanding

1-10 of 10

< Prev Next >

Content Understanding

Microsoft Foundry / demos-resource-project

Search with AI (Ctrl + K)

New Foundry Home Discover Build Operate Docs

Agents

Workflows

Models

Fine-tune

Tools

Knowledge

Data

Evaluations

Guardrails

Azure Content Understanding - Layout

Run analysis

Drag & drop files here or Browse for files

Sample layout_checklist.jpg

Sample layout_financial... rt.jpg

Beacon Roofing Supply, Inc. Notes to Consolidated Financial Statements Year Ended September 30, 2010 (dollars in thousands, except per share data or as otherwise indicated)

5. Prepaid Expenses and Other Current Assets

The significant components of prepaid expenses and other current assets were as follows:

	September 30, 2010	September 30, 2009
Vendor rebates	\$34,538	\$38,857
Refundable income taxes	3,846	3,845
Other	4,731	10,312
	\$43,115	\$52,714

6. Property and Equipment, net

Property and equipment, net, consisted of the following:

	September 30, 2010	September 30, 2009
Land	\$ 3,056	\$ 3,190
Buildings and leasehold improvements	19,987	18,725
Equipment	107,548	100,731
Furniture and fixtures	10,906	10,478
Less: accumulated depreciation and amortization	(141,537)	(133,124)
	\$ 47,751	\$ 52,965

Depreciation and amortization of property and equipment totaled \$17,071, \$17,389 and \$18,381 in 2010, 2009 and 2008, respectively.

7. Accrued Expenses

The significant components of accrued expenses were as follows:

	September 30, 2010	September 30, 2009
Uninvoiced inventory receipts	\$ 7,719	\$19,970
Employee-related accruals	(13,498)	18,736
Unsettled loss on financial derivatives	11,063	17,348
Other	17,811	24,482
	\$50,132	\$75,536

8. Financing Arrangements

Senior Secured Credit Facilities

On November 2, 2008, the Company entered into an amended and restated seven-year \$500 million U.S. senior secured credit facility and a C\$15 million senior secured Canadian credit facility with GE Antares Capital ("GE Antares") and a syndicate of other lenders (combined, the "Credit Facility"). The Credit Facility provides for a cash receipts lock-box arrangement that gives the Company sole control over the funds in lock-box accounts, unless excess availability is less than \$10 million or an event of default occurs, in which case the senior secured lenders would have the right to take control over such funds and to apply such funds to repayment of the senior debt.

Content Result

Markdown Paragraphs Tables

Title

Beacon Roofing Supply, Inc. Notes to Consolidated Financial Statements Year Ended September 30, 2010 (dollars in thousands, except per share data or as otherwise indicated)

SectionHeading

5. Prepaid Expenses and Other Current Assets

Paragraph

The significant components of prepaid expenses and other current assets were as follows:

Paragraph

September 30, 2010

Content Result

JSON

```
1 {
2   "id": "2a933274-e0fd-4c40-85f5-b878bfe931ac",
3   "status": "Succeeded",
4   "result": {
5     "analyzerId": "prebuilt-layout",
6     "apiVersion": "2025-11-01",
7     "createdAt": "2025-11-16T05:41:55Z",
8     "warnings": [],
9     "contents": [
10      {
11        "path": "input1",
12        "markdown": "# Documents to Store and",
13        "kind": "document",
14        "startPageNumber": 1,
15        "endPageNumber": 1,
16        "unit": "pixel",
17        "pages": [
18          {
19            "pageNumber": 1,
20            "angle": 0,
21            "width": 2580,
22            "height": 3433,
```


Content Understanding Extraction Capabilities



Documents

- OCR – printed or handwritten text from 100+ languages
- Content elements (selection mark, check box, radio buttons, barcodes, formulas, image elements)
- Layout elements (pages, paragraphs, lines, sections)
- Tabular data (including complex multi-page layouts)
- Prebuilt Analyzers (including RAG-ready extractions)
- Multi-document input (Pro mode)
- Multi-file reasoning (Pro mode)
- Reference data (Pro mode)
- Classification & splitting (Pro mode)



Image

- Generate titles, descriptions, and summaries
- Classify elements
- Prebuilt Analyzer



Audio

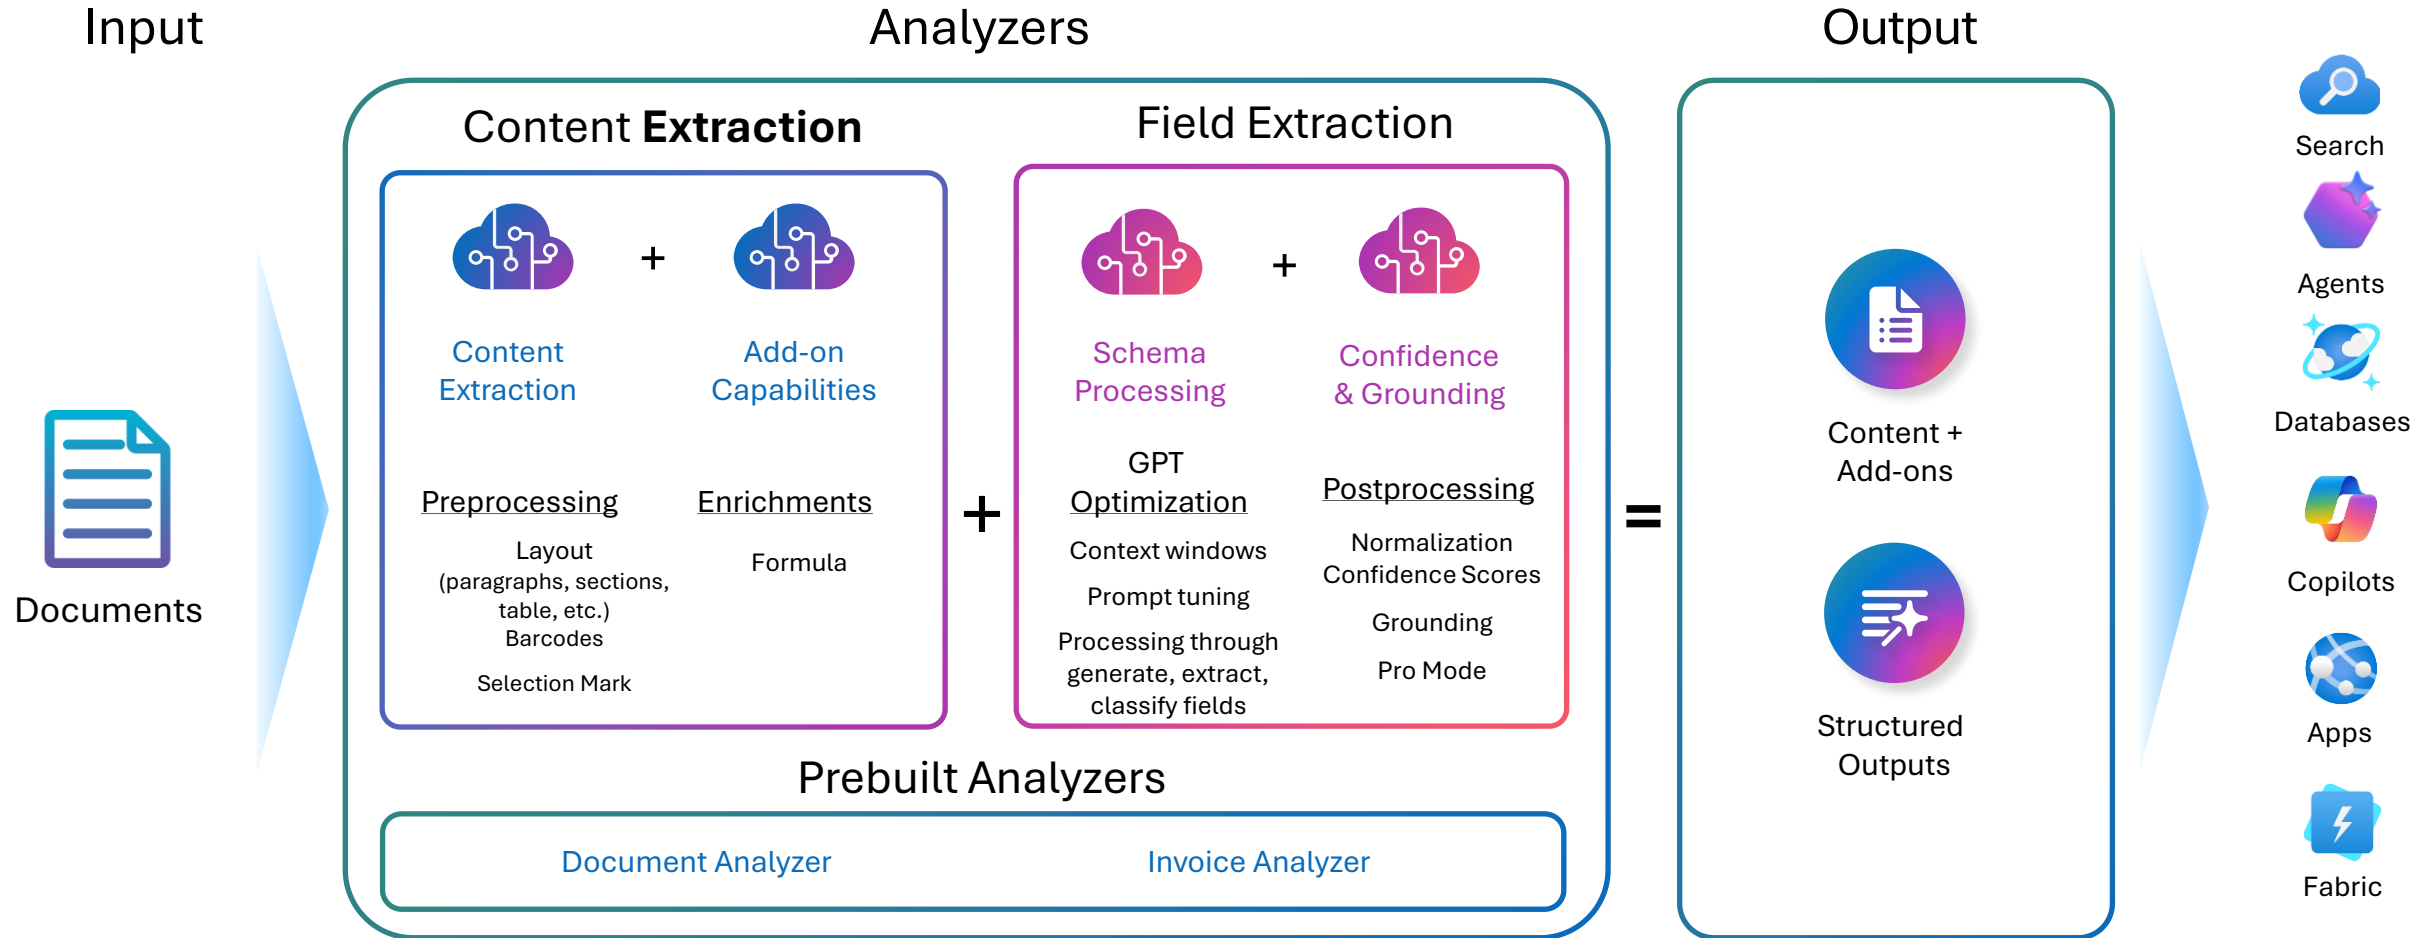
- Transcription
 - Diarization
 - Multi-lingual Transcription
 - Language Detection
 - Multi-Speaker role detection
- Generate conversation summaries and transcripts
- Determine sentiment
- Prebuilt Analyzers



Video

- Transcription
 - Diarization
 - Multi-lingual Transcription
 - Language Detection
 - Key frame extraction
 - Shot detection
- Segmentation (whole-video, automatic, custom)
- Face Detection and Recognition (gated)
- Generate summaries
- Determine sentiment
- Prebuilt Analyzers
 - RAG-ready extraction (prebuilt video indexing)

Document Capabilities



Audio Capabilities







Image Capabilities



Video Capabilities



Pro Analyzers



Agenda

- Présentation Azure Content Understanding
- **Démonstrations**
- Azure Document Intelligence vs Azure Content Understanding
- Ressources et documentation



Démonstrations



Agenda

- Présentation Azure Content Understanding
- Démonstrations
- **Azure Document Intelligence vs Azure Content Understanding**
- Ressources et documentation



Azure Document Intelligence vs Content Understanding

Capability	Document Intelligence	Content Understanding
Focus	Document processing	Generative and reasoning based content understanding
Use cases	Structured/semi structured	Multi modal
Effort	Labeled data for custom models	Zero shot extraction with schema definition
Pricing Model	Per page (fixed rates)	Per page + token-based (variable)
Cost Predictability	★★★★★ — Fixed per-page pricing	★★★★☆ — Varies with schema complexity & model
Document OCR (Read)	☑ Available	☑ Available
Layout Analysis	☑ Available	☑ Available
Prebuilt Invoice/Receipt	☑ Available	☑ Available
Custom Field Extraction	☑ Available	☑ Available
Image Processing	✗ Not supported natively	☑ Free extraction + LLM analysis
Audio Processing	✗ Not supported	☑ \$0.36/min + LLM tokens
Video Processing	✗ Not supported	☑ \$1.00/min + LLM tokens
Multimodal (Doc+Image+Audio+Video)	✗ Documents only	☑ Unified pipeline
LLM Model Choice	✗ No choice	☑ GPT-4o-mini, GPT-4.1, etc.
Confidence Scores	☑ Available	☑ Optional (increases token cost ~2x)
Source Grounding	✗ Limited	☑ Shows where data was extracted
Segmentation / Categorization	✗ Classification only (\$3/1K)	☑ AI-powered segmentation & routing
Reasoning capabilities	✗	☑ Available
Multi file processing	✗	☑ Available
Latency	Low	Medium
Free Tier	☑ 500 pages/month	✗ No free tier
Commitment Tiers	☑ Available for volume discounts	✗ Pay-as-you-go only
Maturity	★★★★★ — GA, battle-tested	★★★★☆ — GA since Nov 2025
Best For	High-volume document-only workflows with predictable pricing	Multimodal pipelines, RAG, agentic workflows, cost-optimized with mini models
Container support	☑ Available	✗ Not available

[Choose the right Azure AI tool for document processing - Foundry Tools | Microsoft Learn](#)

Azure Content understanding

Standard mode vs Pro mode

Feature	Standard mode	Pro mode
 Large documents	✓	✓
 Field mode	✓	✓
 Extract, classify, and generate fields	✓	✓
 Grounding and confidence scores	✓	✗
 Input document type	Documents, images, video, audio	Documents
 Max fields	100	100
 Multiple input document processing	✗	✓
 Reference dataset integration	✗	✓
 Multi-step reasoning	✗	✓

Pro mode is a preview feature available only in the preview API version (2025-05-01-preview).

Cost study 1,000 Pages Scenario

Use Case	Doc Intelligence	Content Understanding	Difference
Basic OCR / Read	\$1.50	~\$1.00	CU is cheaper
Layout Extraction	\$10.00	~\$5.00 extraction (no field extraction)	CU ~50% cheaper for layout-only
Invoice Field Extraction	\$10.00	~\$6.54 (with GPT-4o-mini)	CU ~35% cheaper with mini model
Invoice + Confidence + Source Grounding	\$10.00	~\$11–13 (2x tokens for conf.)	DI slightly cheaper with all features
Custom extraction	\$30.00	~\$6–15 (schema-dependent)	CU potentially much cheaper

Note: Content Understanding costs are estimates based on ~1,100 input + 60 output tokens/page with GPT-4o-mini. Actual costs vary by document complexity.

Agenda

- Présentation Azure Content Understanding
- Démonstrations
- Azure Document Intelligence vs Azure Content Understanding
- **Ressources et documentation**



Documentation

- [Azure Content Understanding documentation | Microsoft Learn](#)
- [Quickstart: Azure Content Understanding in Foundry Tools REST APIs - Foundry Tools | Microsoft Learn](#)
- [Service quotas and limits - Content Understanding - Foundry Tools | Microsoft Learn](#)
- [Content Understanding region and language support - Foundry Tools | Microsoft Learn](#)
- [Pricing - Azure Content Understanding in Foundry Tools | Microsoft Azure](#)
- <https://github.com/Azure-Samples/azure-ai-content-understanding-with-fabric>

Multimodal data processing with Azure AI Content Understanding

MULTIMODALITY USING AZURE AI

Azure AI Content Understanding: How to Get Started

with Vinod Kurpad

07:41

EPISODE

Azure AI Content Understanding: How to Get Started

Apr 1, 2025

Azure

MULTIMODALITY USING AZURE AI

Post-call analytics using Azure AI Content Understanding

with Jan Görgen

06:59

EPISODE

Post-call analytics using Azure AI Content Understanding

Apr 1, 2025

Azure

MULTIMODALITY USING AZURE AI

RAG on Video using Azure AI Content Understanding

with Joe Filcik

06:20

EPISODE

RAG on Video using Azure AI Content Understanding

Apr 1, 2025

Azure

Article Medium et GitHub Repo



<https://aka.ms/azure-cu-article>



<https://aka.ms/azure-cu-repo>

Demos videos

The screenshot displays the Azure Content Understanding web interface. At the top, there's a YouTube logo and a search bar labeled 'Rechercher'. The main header shows 'Azure AI Foundry' and 'Content Understanding'. A sidebar on the left lists various features like Overview, Model catalog, Playgrounds, AI Services, Build and customize, Agents, Templates, Fine-tuning, Content Understanding (selected), Prompt flow, Observe and optimize, Tracing, Monitoring, Protect and govern, Evaluation, and Guardrails + controls.

In the center, a 'Create a new task' dialog is open. It has fields for 'Task name' (containing 'Invoice analysis') and 'Description' (containing 'Analyze the'). Below these are two options: 'Single-file content analysis' (marked as 'Recommended') and 'Multi-file content analysis' (marked as 'Preview'). The 'Single-file content analysis' option includes a description: 'Extract structured output from unstructured data using a custom-defined schema. Recommended for: most use cases building content analyzer.' The 'Multi-file content analysis' option includes a description: 'Generate insights across multiple files using contextual reference knowledge. Recommended for: complex use cases require reasoning processing.'

To the right of the dialog, there's a table with columns 'Created On' and a trash icon. It lists several tasks with dates ranging from 7/22/2025 to 6/10/2025.

On the far right, a video player interface is shown with the title 'Azure Content Understanding' and 'Serge Analytics - 1/5'. Below the title is a list of five demo videos:

- Azure Content Understanding Document Processing (2:16)
- Azure Content Understanding Image Processing (1:45)
- Azure Content Understanding Pro mode (7:18)
- Azure Content Understanding Video Processing (2:06)
- Azure Content Understanding Audio Processing (3:13)

[Azure Content Understanding - YouTube](#)

This deck is available here:

Prochain évènement



Présentation des accélérateurs Citadel et Spec2Cloud

Citadel et Spec2Cloud sont deux accélérateurs Microsoft IA. Citadel fournit une **gouvernance centralisée, la conformité et l'observabilité pour déployer des agents IA sécurisés et contrôlés à l'échelle**. Spec2Cloud accélère le **développement d'applications et d'agents IA via des modèles de spécifications**.

Ensemble, ils résolvent le dilemme vitesse vs contrôle : Spec2Cloud vous fait développer vite, Citadel vous assure de le faire en sécurité et conformité. Ces solutions offrent une base pour **transformer rapidement les idées en agents IA gouvernés et production-ready**.



Serge RETKOWSKY

France AI Global Black Belt
@Microsoft



Alibek JAKUPOV

France AI Global Black Belt
@Microsoft



Nicolas ROBERT

Profession
@Cellenza
AI MVP

LinkedIn Live

jeudi

19/03/2026

à partir de 13h

https://lnkd.in/eUe_bBrz

<https://www.linkedin.com/company/global-ai-community-france/posts/?feedView=all>

Replays Global AI Community



[\(7\) Global AI Community Virtuelle Francophone: Posts | LinkedIn](#)

Merci